

# Railway Scheduling Using Reinforcement Learning

*A Project Report Submitted  
in Partial Fulfillment of the Requirements  
for the Degree of*

**Bachelor of Technology**

*by*

**Arpit Singh**  
(111601031)

*under the guidance of*

**Dr. Chandra Shekar Lakshminarayanan**



INDIAN INSTITUTE  
OF TECHNOLOGY  
**PALAKKAD**

**DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING**

# Contents

<b>List of Figures</b>	<b>2</b>
<b>1 Introduction</b>	<b>1</b>
<b>2 Problem Statement</b>	<b>1</b>
2.1 Indian Railways . . . . .	1
2.2 Railway Scheduling Problem . . . . .	2
2.2.1 Scheduling . . . . .	2
2.2.2 Rescheduling . . . . .	3
<b>3 Reinforcement Learning Approach</b>	<b>4</b>
3.1 Brief Description . . . . .	4
3.2 Prior Work . . . . .	4
3.2.1 State Representation (Local state space) . . . . .	4
3.2.2 Action and Policy Definition . . . . .	5
3.3 Proposed Approach . . . . .	5
3.3.1 State Space Representation . . . . .	5
3.3.2 Action And Policy Definition . . . . .	6
3.4 Objective Function . . . . .	6
<b>4 Implementation</b>	<b>7</b>
4.1 Railway simulator . . . . .	7
4.1.1 Requirement . . . . .	7
4.1.2 Implementation . . . . .	7
<b>5 Conclusion and Future Work</b>	<b>9</b>
<b>References</b>	<b>10</b>

# List of Figures

2.1	The line and junction topology of railway networks in India [1]. . . . .	1
2.2	Linear Railway Lines [2]. . . . .	2
3.1	Mapping train location and direction of movement to resource status, relative to the ‘current train’ [2]. . . . .	5

# Introduction

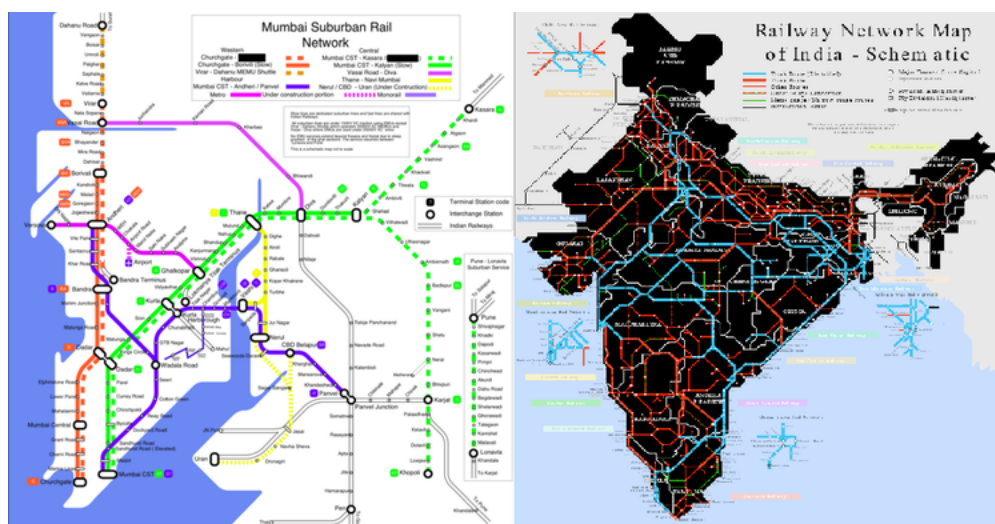
The aim is to work on an algorithm for scheduling bidirectional railway lines (both single- and multi-track) using a framework of Reinforcement Learning. Given deterministic arrival/departure times for all the trains on the lines, their initial positions, priority and halt times, traversal times, deciding on track allocations is a job shop scheduling problem (NP Complete ). However, due to the stochastic nature of the delays, the track allocation decisions have to be made in a dynamic manner, while minimising the total priority-weighted delay. This makes the underlying problem one of decision making in of stochastic event driven systems. The primary advantage of the proposed algorithm compared to exact approaches is its scalability, and compared to heuristic approaches is its solution quality. Improved solution quality is obtained because of the inherent adaptability of reinforcement learning to specific problem instances.

This report presents the problem statement, discusses about the approach and future scopes of improvement.

# Problem Statement


## Indian Railways

Let us first describe the nature of the railway system in this country. The Indian railway network is designed to consist of long ‘lines’ (a string of stations), which connect with each other at ‘junction’ stations. Each station is composed of one or more parallel **tracks**, which may be associated with a fixed direction of traffic, or they could be bidirectional. Similarly, there are one or more tracks between each neighbouring pair of stations. These tracks are typically referred to as **sections**, in order to differentiate them from tracks actually at a station. The section tracks can also be unidirectional (fixed direction of train movement) or bidirectional. The Indian network typically consists of sections with one or two tracks.



**Fig. 2.1** The line and junction topology of railway networks in India [1].

The approach we are focussing on now deals with linear railway networks with multiple parallel tracks, of the type shown in figure 2.2. This restriction on topology is reasonable because rail networks are designed in the form of multi-station linear arcs connected at junction stations.



**Fig. 2.2** Linear Railway Lines [2].

## Railway Scheduling Problem

### Scheduling

A specific problem instance begins by defining the resources on the railway line, as given by the number of stations, their order, and the number of parallel tracks (both at each station and between two neighbouring stations). Besides resource level information, train movements over the scheduling horizon must be described in one of two ways.

- To define a reference timetable which gives the desired arrival and departure time of each train at each station.
- To provide the earliest movement times from their current locations (or origin stations), followed by the minimum running times (on track sections between stations) and halt times (at stations) up to the destinations.

Note that the running and halt times can be completely heterogeneous: each train may have a different running/halt time in each resource, depending on the length of track, the type of halt, and the type of locomotive.

Timetabling refers to an offline planning problem for a railway network. **Given a set of trains and their origins and destinations (with or without a fixed route), the goal is to assign track resources for each train for a fixed time period, such that they all complete their journeys without conflicts.**

Such a timetable may be infeasible if the desired arrival and departure times violate the track usage rules defined earlier. The task of the scheduling algorithm is to adjust the arrival and departure times such that all rules are satisfied, while minimizing an objective called priority-weighted delay. This schedule is to be computed for all trains up to their destinations.

The railway problem has been shown in literature to be a **‘blocking’ version of the Job Shop Scheduling Problem (JSSP)**, where the job (train) must wait in the current resource (track) until the next resource is freed (there is no buffer for storing jobs between two resources). This version of the JSSP is also **NP complete**, with the result that exact solutions require an exponential amount of time for computation.

## Rescheduling

Another problem is that of rescheduling. Rescheduling is the online counterpart of the timetabling problem, where the goal is to recover from a disruption to the timetable, caused by delays or faults. The mathematical differences are found in two aspects.

- The goal is to return to the original timetable using built-in slack times, instead of defining the timetable itself. This implies that the objective function would focus on minimizing delays to trains with respect to the timetable, or the time required for deviations to be smoothed out.
- The online nature of the problem implies that there is very limited time available to compute solutions, and that sub-optimal but reasonably efficient solutions are acceptable.

# Reinforcement Learning Approach

## Brief Description

Reinforcement learning works by learning to map the state of the system to the choice of available actions based on long-term reward (or penalty). In this report we will discuss review in brief the prior work [2] and then discuss the work that will be done as part of the project. The key difference in both the the approaches is:

- In the prior approach, we find state space with respect to train and it is the train that is making the decision (So we have the environment where each train is acting as the agent i.e. multiagent environment)
- In the proposed work, we have the central controller as the single agent and it is that agent that is scheduling the trains.

## Prior Work

### State Representation (Local state space)


Here, we compute the state space as a function of **local neighborhood** of each train. A state vector is computed for each train every time a decision about its next move is to be computed. Relative to the direction of motion, we define resources as being behind (in the direction opposite to the direction of motion) or in front (in the direction of motion) of the train. A user-defined finite number of resources  $l_b$  behind each train and  $l_f$  in front of each train are used for defining the state vector. These are referred to as local resources. Including a few resources behind the train in the state definition ensures that overtaking opportunities for fast-moving trains are not missed. The total number of local resources is  $(l_b + 1 + l_f)$ .

The entry in the state vector corresponding to each local resource takes one of  $R$  integer values  $\{0, 1, 2, \dots, R - 1\}$ , referred to as the status  $S_r$  of resource  $r$ . Higher values indicate higher congestion within the resource, and are driven by the number of occupied tracks. Let us define the number of tracks in resource  $r$  to be equal to  $N_r$ , out of which  $T_{r,c}$  tracks contain trains converging with (heading towards) the current train, while  $T_{r,d}$  tracks contain trains diverging from (heading away from) the current train. Since at most one train can occupy a given track, we note that  $T_{r,c} + T_{r,d} < N_r$ . The mapping from track occupancy to resource status is,

$$S_r = R - 1 - \min(R - 1, \lceil N_r - w_c T_{r,c} - w_d T_{r,d} \rceil).$$

Here,  $0 \leq w_c, w_d \leq 1$  are weights that can de-emphasise the effect of converging and diverging trains on the perceived status of a resource. We have one more entry in the state space for the priority of trains. If we assume that the model accommodates up to  $P$  priority





**Fig. 3.1** Mapping train location and direction of movement to resource status, relative to the ‘current train’ [2].

levels, the size of the state space is equal to  $P \cdot R^{l_b+1+l_f}$ . Note that this value does not depend on the scale of the problem instance, in terms of the number of trains, the lengths of their journeys, and the number of resources.

### Action and Policy Definition

In this approach, state space is computed for each train locally. The choice of actions in any given state is binary, with 0 representing a decision to move the current train to the next resource on its journey, and 1 representing a decision to halt in the current resource for a predefined time period (1 minute in this paper). If the train is halted, the decision-making procedure is repeated after the time period elapses. The policy that can be used to take the action is  $\epsilon$ -greedy with the value of  $\epsilon$ -decreasing over the training. It comprise of three steps

- Select the train on which to take the action.
- Compute the state space for the train.
- Choose the action depending on the  $\epsilon$ -greedy policy.

## Proposed Approach

### State Space Representation

We are planning to use the whole network along with the positions of each train to take as the state space. The key idea is to let the RL algorithm find which part of the state space is important to make decision. In the prior approach, we are kind of tuning how far the train can see and it’s action depends only on local neighborhood but in our work we are going to

include everything in the state space. Since, including everything can blow the size of the state space really quickly, so to tackle with this problem we can use function approximation methods (Deep Q-Learning) to learn the state space and then work on it.

### Action And Policy Definition

In our proposed approach, we take the whole network topology and the position of the trains as the state space. So we have the central controller and that central controller will take the action for each train depending on state space that takes everything into account. At a particular time during the simulation, let's say we have n number of trains waiting for the action (ready to move or stop), then the controller have the action space of  $2^n$  and it has to make decision for each of the trains (although in some order, as we are running the simulator and taking action for one train at a time). Size of the action space depends on the number of trains waiting for the action to be taken. Here again we can use the  $\epsilon$ -greedy policy for exploration in the initial phase and then further exploitation.

### Objective Function

A number of objective functions have been used in the railway scheduling context, in order to achieve goals such as delay reduction, passenger convenience, and timetable robustness. One of the commonly used measures of schedule quality is priority-weighted delay. A delay is defined to be the non-negative difference between the time of an event as computed by the algorithm, and the desired time as specified by the timetable. The priority-weighted average delay is the mean over all trains and all stations of individual delays divided by train priorities. This quantity is used as the objective function, but the algorithm can accommodate other measures equally easily (for example, a non-linear function of delays in order to increase fairness of delay distribution).

$$J = \frac{1}{N_{r,t}} \sum_{r,t} \frac{\delta_{r,t}}{P_t}$$

where  $\delta_{r,t}$  is the delay for train t on departure from resource r,  $p_t$  is the priority of train t, and  $N_{r,t}$  is the total number of departures in the schedule. Note that this expression includes all events for all trains, for their entire journey.

Prior work uses priority weighted delay as the objective function. In our proposed approach as well, we can shape the reward functions as to use the same objective function.

# Implementation

As for now, we are focussing on how to implement the first approach and then move onto the second approach. The integrated reinforcement learning algorithm is driven by a discrete event simulator. There are already some railway simulators like **OpenTrack** [3] and **RailML**[4] but they would be useful for the final analysis of the results. Once we have the desired timetable then we can use these simulator softwares to determine the quality of solution. But for the implementation of the algorithm we have to implement the simulator on own.

## Railway simulator

### Requirement

The simulator is supposed to be robust enough that it can run both toy and real life examples. The simulator is suppose to run through several episodes during training and hence need to be efficient. At the beginning of every episode, the initial locations of all the trains are reset to their original values. It is assumed that trains that have not yet started, or have finished their journeys, do not occupy any of the tracks. Following the train-to-resource mapping, the simulator creates a list of events for processing, one corresponding to each train (whether already running or yet to start its journey). Each event in the list contains the following information: the time at which to process the event, the train to which it corresponds, the resource where the train is currently located, the last observed state-action pair for the train (empty if the train is yet to start), and the direction of the train journey.

At each step, the algorithm moves the simulation clock to the earliest time stamp in the event list. If multiple events are to be processed at the same time stamp, they are handled sequentially. We are for now not focussing on how to avoid deadlock, but instead if we get into deadlock, we will detect and give huge negative reward and the RL algorithm is suppose to avoid deadlock on it's own.

### Implementation

There are two components to railway simulator :

1. Underlying Railway Network.
2. Trains and the simulation of there movements.

For the implementation of the railway network we can use **NetworkX**[5] package of python. **NetworkX is a Python package for the creation, manipulation, and study of the structure, dynamics, and functions of complex networks.**

Once the network is ready we have to simulate movement of each train over the network. For that we can use **SimPy**[6] package of python. **SimPy is a process-based discrete-event simulation framework based on standard Python.** In this we can model each train as the separate process and network as the resource. We can model the movement of trains using this package. We yield events when the train starts from some station and once the train reaches the next station, event is yielded and then we can process accordingly. So whole simulation is done by generating events at points where the algorithm is supposed to take action.

# Conclusion and Future Work

So far, focus is to fully understand the problem statement, what are the variations of the problem and then how RL algorithms can be used to solve this problem. The report discusses two approach, to solve the problem. Currently the focus is to implement the prior approach and see how good it is working, how good it is scaling to real life problem instances, how good it is performing compared to the present approaches and how to improve upon it. For the implementation of the algorithm, we need to implement the discrete railway simulator. For that, I have gone through the NetworkX and SimPy packages of standard python. The future plan is to complete the implementation of the simulator and then test our proposed method on the simulator.

Since we are tackling blocking version of the JSSP problem, so the approach that we will develop can be used to solve the JSSP problem with reasonable approximation. So the future plan is also to use the developed approaches on other similar problems as well.

# References

- [1] H. KHADILKAR, “Artificial intelligence for indian railways,” <https://www.ai4bharat.org/articles/railways>, June 2019.
- [2] H. Khadilkar, “A scalable reinforcement learning algorithm for scheduling railway lines,” *IEEE*, 2018.
- [3] [http://www.opentrack.ch/opentrack/opentrack\\_e/opentrack\\_e.html](http://www.opentrack.ch/opentrack/opentrack_e/opentrack_e.html).
- [4] <https://www.railml.org/en/>.
- [5] <https://networkx.github.io/>.
- [6] <https://simpy.readthedocs.io/en/latest/contents.html>.